

05: MAXIMUM LIKELIHOOD ESTIMATION

Stat250-S25

1 WARM UP

Example: You are testing seeds from a new plant variety. You plant 10 seeds ($n=10$) and observe that 3 of them successfully germinate ($k=3$). Consider two hypotheses about the true germination probability (p) for this variety:

Hypothesis A:

Hypothesis B:

Given you observed 3 germinations out of 10 seeds, which hypothesis (A or B) is more supported by the data?

In general, the *likelihood* of seeing $X = 3$ given p is:

3 ways of finding the maximum:

1. Approximate solution graphically

```
germ_function <- function(p) {choose(10,3) * p^3 * (1-p)^7} # define function
ggplot() +
  geom_function(fun = germ_function) +
  xlim(c(0,1))
```



2. Find a numerical approximation

```
optimize(germ_function, interval = c(0,1), maximum = TRUE)
```

```
$maximum  
[1] 0.3000157
```

```
$objective  
[1] 0.2668279
```

3. Use calculus to find an exact maximum

Can we generalize to *any* data?

2 DEFINITIONS

Likelihood function

Let $f(x; \theta)$ denote the probability mass function for a discrete distribution with associated parameter θ . Suppose X_1, X_2, \dots, X_n are a random sample from this distribution and x_1, x_2, \dots, x_n are the actual observed values. Then, the *likelihood function* of θ is:

Maximum likelihood estimate

A maximum likelihood estimate (MLE), $\hat{\theta}_{MLE}$ is the value of θ that maximizes the likelihood function, or equivalently, that maximizes the log-likelihood $l(\theta) = \ln L(\theta)$

Bernoulli Data Example

Estimator

Estimate

3 EXERCISE

Let X_1, \dots, X_n be an iid random sample from a distribution with PDF

$$f(x|\theta) = (\theta + 1)x^\theta, 0 \leq x \leq 1$$

1. Find the maximum likelihood estimator for a random sample of size n using calculus.
2. Suppose we observe a sample of size 5: $\{.83, .49, .72, .57, .66\}$. Find the maximum likelihood estimate and verify with a graph or numerical approximation